

Filippo Ficarra: Lab 2 - Feature extraction and Optical flow

fficarra@student.ethz.ch, 22-938-062.

09/10/2023 - 01:48h

1 Introduction

In this lab we aim to extract features from images, such as corners, using Harris detection and match descriptors between two different photos of the same thing.

2 Harris corner detection

The idea of Harris corner detection is to analyze the change of intensity of pixels in a window. Using this idea we can basically identify three different regions:

- flat: there is no change of intensity in all the directions
- edge: there is no change of intensity in the direction of the edge
- corner: large change of intensity in all the directions

To perform this analysis we define a window \mathbf{W} and we define the SSD error as:

$$E(u, v) = \sum_{(x, y) \in \mathbf{W}} [I(x + u, y + v) - I(x, y)]^2$$

using Taylor approximation we can approximate the error with:

$$E(\Delta x, \Delta y) \approx [\Delta x, \Delta y] M \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}, M = \sum \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Note that I_x and I_y are respectively the partial derivative of the image with respect to x and y .

2.1 Image derivatives

Since images are discrete we need to compute some approximation of the derivative. The approximation used in this case is the following:

$$I_x = \frac{I(x+1, y) - I(x-1, y)}{2}, I_y = \frac{I(x, y+1) - I(x, y-1)}{2}$$

We can exploit the convolution operations to compute these derivatives, using the following filters:

$$\text{filter}_x = \begin{bmatrix} -0.5 & 0 & 0.5 \end{bmatrix}, \text{filter}_y = \begin{bmatrix} -0.5 \\ 0 \\ 0.5 \end{bmatrix}$$

This convolutions have been done in python in the following way:

Listing 1: Image gradients

```

1  filter_x = np.array([[1/2, 0, -1/2]])
2  filter_y = np.array([[1/2], [0], [-1/2]])
3
4  I_x = signal.convolve2d(img, filter_x, mode='same')
5  I_y = signal.convolve2d(img, filter_y, mode='same')

```

2.2 Gaussian blur and local-auto correlation matrix

The next step is to introduce blur in the image to make the detection more robust. Applying a Gaussian Blur filter prior to performing corner detection serves the purpose of decreasing image noise, thereby enhancing the outcome of corner-detection.

To do so we use a gaussian kernel g , thus the matrix M is the following:

$$M = g * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Furthermore to compute the response function C we will need just this 3 elements:

Listing 2: Local auto-correlation matrix elements

```

1  I_xx_blur = cv2.GaussianBlur(I_x**2,(5,5),sigma,borderType=cv2.BORDER_REPLICATE)
2
3  I_yy_blur = cv2.GaussianBlur(I_y**2,(5,5),sigma,borderType=cv2.BORDER_REPLICATE)
4
5  I_xy_blur = cv2.GaussianBlur(I_x*I_y,(5,5),sigma,borderType=cv2.BORDER_REPLICATE)
6
7  # 3. Compute elements of the local auto-correlation matrix "M"
8
9  g_xx = I_xx_blur
10 g_yy = I_yy_blur
11 g_xy = I_xy_blur

```

2.3 Harris response and corner detection

The Harris detector uses the following function to score the presence of corners:

$$C = \lambda_1 \lambda_2 - k(\lambda_1 + \lambda_2)^2 = \det(M) - k(\text{trace}(M))^2$$

therefore if

- $\lambda_1 \sim 0$ and $\lambda_2 \sim 0 \implies C \ll 0$ and the region is flat, since the intensity of the pixels doesn't really change in that region
- $\lambda_2 \gg \lambda_1 \implies C < 0$ and we detect an horizontal edge
- $\lambda_1 \gg \lambda_2 \implies C < 0$ and we detect a vertical edge
- $\lambda_1 \sim \lambda_2$ and both large, $\implies C > 0$ and we detect a corner

Alternatively to use $\det(M)$ and $\text{trace}(M)$, we can rewrite the response function like this:

$$C = g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - k[g(I_x^2) + g(I_y^2)]^2$$

and then we can reuse the components derived above:

Listing 3: Harris response function

```

1  C = (g_xx * g_yy - (g_xy**2)) - k * (g_xx + g_yy)**2

```

In order to detect corners we need just to compare every entry of the response with a threshold and perform non-maximum suppression. A non-maximum suppression process allows selecting a unique feature in each neighborhood. This has been performed in python like this:

Listing 4: corners

```
1 corners = np.argwhere((C > thresh) & (C == ndimage.maximum_filter(C, (3,3))))
2 corners = corners[:, [1, 0]] # convention row = y, column = x
```

2.4 Results

In this section I'll show the result of the detection for the different parameters Sigma (Gaussian Blur), k (Response function) and the threshold.

As expected, when one decreases the threshold, the number of corners detected is greater but not so accurate. Increasing the value of k reduces the number of points since the response function gets smaller as shown in the first two pairs of images (here every parameter is fixed and $k = 0.04$ and $k = 0.06$). Increasing k we will detect less edges, but also real corners. Furthermore also increasing the sigma should reduce the number of points, as shown in the house image, since the image has been smoothed. It can happen though that the noise level is already high in the image and therefore increasing the blurring can help to find more corners, as shown in the blocks image below ($\sigma = 0.5$ and $\sigma = 2.0$).




Figure 1: $k = 0.04$, $\sigma = 0.5$, $t = 1e-4$ and keypoints = 5




Figure 2: $k = 0.04$, $\sigma = 0.5$, $t = 5e-5$ and keypoints = 13




Figure 3: $k = 0.04$, $\sigma = 0.5$, $t = 1e-5$ and keypoints = 34




Figure 4: $k = 0.04$, $\sigma = 0.5$, $t = 5e-6$ and keypoints = 46




Figure 5: $k = 0.04$, $\sigma = 0.5$, $t = 1e-6$ and keypoints = 127



Figure 6: $k = 0.06$, $\sigma = 0.5$, $t = 1e-4$ and keypoints = 4




Figure 7: $k = 0.06$, $\sigma = 0.5$, $t = 5e-5$ and keypoints = 10




Figure 8: $k = 0.06$, $\sigma = 0.5$, $t = 1e-5$ and keypoints = 28




Figure 9: $k = 0.06$, $\sigma = 0.5$, $t = 5e-6$ and keypoints = 37




Figure 10: $k = 0.06$, $\sigma = 0.5$, $t = 1e-6$ and keypoints = 86




Figure 11: $k = 0.04$, $\sigma = 0.5$, $t = 1e-4$ and keypoints = 23




Figure 12: $k = 0.04$, $\sigma = 0.5$, $t = 5e-5$ and keypoints = 51




Figure 13: $k = 0.04$, $\sigma = 0.5$, $t = 1e-5$ and keypoints = 134




Figure 14: $k = 0.04$, $\sigma = 0.5$, $t = 5e-6$ and keypoints = 173



Figure 15: $k = 0.04$, $\sigma = 0.5$, $t = 1e-6$ and keypoints = 222




Figure 16: $k = 0.06$, $\sigma = 0.5$, $t = 1e-4$ and keypoints = 19




Figure 17: $k = 0.06$, $\sigma = 0.5$, $t = 5e-5$ and keypoints = 42




Figure 18: $k = 0.06$, $\sigma = 0.5$, $t = 1e-5$ and keypoints = 115




Figure 19: $k = 0.06$, $\sigma = 0.5$, $t = 5e-6$ and keypoints = 154



Figure 20: $k = 0.06$, $\sigma = 0.5$, $t = 1e-6$ and keypoints = 191




Figure 21: $k = 0.06$, $\sigma = 0.5$, $t = 1e-4$ and keypoints = 4




Figure 22: $k = 0.06$, $\sigma = 0.5$, $t = 5e-5$ and keypoints = 10




Figure 23: $k = 0.06$, $\sigma = 0.5$, $t = 1e-5$ and keypoints = 28




Figure 24: $k = 0.06$, $\sigma = 0.5$, $t = 5e-6$ and keypoints = 37




Figure 25: $k = 0.06$, $\sigma = 0.5$, $t = 1e-6$ and keypoints = 86




Figure 26: $k = 0.06$, $\sigma = 2.0$, $t = 1e-4$ and keypoints = 8




Figure 27: $k = 0.06$, $\sigma = 2.0$, $t = 5e-5$ and keypoints = 18




Figure 28: $k = 0.06$, $\sigma = 2.0$, $t = 1e-5$ and keypoints = 41




Figure 29: $k = 0.06$, $\sigma = 2.0$, $t = 5e-6$ and keypoints = 50




Figure 30: $k = 0.06$, $\sigma = 2.0$, $t = 1e-6$ and keypoints = 79




Figure 31: $k = 0.06$, $\sigma = 0.5$, $t = 1e-4$ and keypoints = 19




Figure 32: $k = 0.06$, $\sigma = 0.5$, $t = 5e-5$ and keypoints = 42




Figure 33: $k = 0.06$, $\sigma = 0.5$, $t = 1e-5$ and keypoints = 115




Figure 34: $k = 0.06$, $\sigma = 0.5$, $t = 5e-6$ and keypoints = 154




Figure 35: $k = 0.06$, $\sigma = 0.5$, $t = 1e-6$ and keypoints = 262




Figure 36: $k = 0.06$, $\sigma = 2.0$, $t = 1e-4$ and keypoints = 26




Figure 37: $k = 0.06$, $\sigma = 2.0$, $t = 5e-5$ and keypoints = 53




Figure 38: $k = 0.06$, $\sigma = 2.0$, $t = 1e-5$ and keypoints = 103




Figure 39: $k = 0.06$, $\sigma = 2.0$, $t = 5e-6$ and keypoints = 115




Figure 40: $k = 0.06$, $\sigma = 2.0$, $t = 1e-6$ and keypoints = 144

3 Match descriptors

The evaluation was done with these parameters

Listing 5: parameters used to match descriptors

```
1 HARRIS_SIGMA = 1.0
2 HARRIS_K = 0.05
3 HARRIS_THRESH = 1e-5
```

3.1 Filter keypoints

The function to extract descriptors was already given, we needed to filter the keypoints found with the method above. Filtering in this case means to remove the points on the border of the image. The code is the following

Listing 6: filter keypoints

```
1 def filter_keypoints(img, keypoints, patch_size = 9):
2     '''
3     Inputs:
4     - keypoints:      (q, 2) numpy array of keypoint locations [x, y]
5     Returns:
6     - keypoints:      (q', 2) numpy array of keypoint locations [x, y]
7                       that are far enough from edges
8     '''
9     keypoints = keypoints[(keypoints[:, 0] > patch_size // 2)
10                          & (keypoints[:, 0] < img.shape[1] - patch_size // 2)
11                          & (keypoints[:, 1] > patch_size // 2)
12                          & (keypoints[:, 1] < img.shape[0] - patch_size // 2)]
13
14     return keypoints
```

3.2 SSD

Next, to match descriptors, we need to compute the distances between them. The smaller the distance, the more correlated the keypoints in two images are. This can be done with `cdist` function in `scipy.spatial.distance`.

Listing 7: SSD

```
1 def ssd(desc1, desc2):
2     '''
3     Sum of squared differences
4     Inputs:
5     - desc1:          - (q1, feature_dim) descriptor for the first image
6     - desc2:          - (q2, feature_dim) descriptor for the first image
7     Returns:
8     - distances:      - (q1, q2) numpy array storing the squared distance
9     '''
10    assert desc1.shape[1] == desc2.shape[1]
11    return cdist(desc1, desc2, 'sqeuclidean')
```

The output of the function is for every descriptor in the first image, the distance with every descriptor in the second image

$$dist[i, j] = \sum_{k=1}^{\text{feature_dim}} (desc1[i, k] - desc2[j, k])^2$$

3.3 One way matching

One way matching aims to find for each descriptor in the first image, the correspondent descriptor with lower distance in the second image.

Listing 8: One way matching

```
1  if method == "one_way": # Query the nearest neighbor for each keypoint in image 1
2      matches = np.argmin(distances, axis=1)
3      matches = np.array([(i, match) for i, match in enumerate(matches)])
```

The following image show the result




Figure 41: One way matching

As the figure shows, the one way is not very accurate since it catch some matching that are not correct (this is shown by the crossing connections).

3.4 Mutual matching