

# Filippo Ficarra: Lab 4 - Object Recognition

fficarra@student.ethz.ch, 22-938-062.

04/11/2023 - 18:21h

## 1 Introduction

## 2 Bag of Words

### 2.1 Grid

## 3 CNN-based Classifier

### 3.1 Training

The model is a simplified version of the VGG architecture. The architecture given was the following:

Block Name	Layers	Output Size
conv_block1	ConvReLU (k=3) + MaxPool2d(k=2)	[bs, 64, 16, 16]
conv_block2	ConvReLU (k=3) + MaxPool2d(k=2)	[bs, 128, 8, 8]
conv_block3	ConvReLU (k=3) + MaxPool2d(k=2)	[bs, 256, 4, 4]
conv_block4	ConvReLU (k=3) + MaxPool2d(k=2)	[bs, 512, 2, 2]
conv_block5	ConvReLU (k=3) + MaxPool2d(k=2)	[bs, 512, 1, 1]
classifier	Linear+ReLU+Dropout+Linear	[bs, 10]

Figure 1: VGG simplified architecture

The model was trained with the following parameters:

Argument	Value	Description
batch_size	128	Batch size
log_step	100	How many steps to log once
val_step	100	Validation step
num_epoch	50	Maximum number of training epochs
fc_layer	512	Number of features in the first linear layer in VGG
lr	0.0001	Learning rate

As shown in the table above the model was trained for 50 epochs with batch size 128.

Furthermore the training has been performed on Euler with 1 gpu RTX 4090.

Below we can find the graphs for the train losses in which we can see that the loss progressively drops from an initial value of 2 to a value around 1.3, while the validation accuracies pass from around 45% to a final value of 81.92%




Figure 2: Train Losses



Figure 3: Validation Accuracies

## 3.2 Testing

The model after training with 50 epochs achieved a test accuracy of 80.86 on the CIFAR10 dataset. The test script was run locally on cpu for convenience and gave the following output:

```
79it [00:08, 9.11it/s]
test accuracy: 80.86
```